

Information Retrieval in the Medical Domain: A Comparative Study on TREC-COVID

Lounès KEBDI
Université Paris-Saclay
Gif-sur-Yvette, France
lounes.kebdi@universite-paris-saclay.fr

Lubin LONGUEPEE
Université Paris-Saclay
Gif-sur-Yvette, France
lubin.longuepee@universite-paris-saclay.fr

Raphael LEONARDI
Université Paris-Saclay
Gif-sur-Yvette, France
raphael.leonardi@universite-paris-saclay.fr

Abstract

The COVID-19 pandemic highlighted the critical need for efficient Information Retrieval (IR) systems capable of navigating rapidly expanding scientific literature. A major challenge in this domain is the "vocabulary mismatch" problem, where queries in natural language do not match the technical jargon used in medical papers. This report presents a comprehensive comparative analysis of various IR models on the TREC-COVID dataset, ranging from classical lexical approaches (TF-IDF) to state-of-the-art neural models. We evaluate the performance of Word2Vec (general and biomedical), Deep Averaging Networks (DAN), and Transformer-based models (BioBERT, Sentence-BERT). Our experiments demonstrate that while generic embeddings provide minor improvements over baselines, domain-specific semantic models like Sentence-BERT achieve superior performance (NDCG@10 of 0.8171), significantly outperforming standard BERT implementations. We provide detailed analysis of implementation challenges, including overfitting issues with trainable models and the critical importance of task-specific fine-tuning for transformer architectures.

CCS Concepts

• **Information systems** → **Information retrieval**; • **Computing methodologies** → **Natural language processing**; *Neural networks*.

Keywords

Information Retrieval, TREC-COVID, NLP, Sentence-BERT, Word Embeddings, Biomedical Text Mining

1 Introduction

The rapid proliferation of scientific literature during the COVID-19 pandemic created an urgent need for specialized Information Retrieval (IR) systems. The TREC-COVID challenge [10] was established to evaluate the ability of IR systems to retrieve relevant biomedical articles from a corpus of over 171,000 scientific papers.

The core problem addressed in this study is the *vocabulary mismatch*. Users often query for symptoms using lay terms (e.g., "loss of smell"), while scientific documents use precise medical terminology (e.g., "anosmia"). Traditional exact-match systems (like TF-IDF) fail to capture these semantic equivalences, leading to poor retrieval performance.

Figure 1 illustrates this vocabulary mismatch problem and how semantic models address it.

This report presents a comprehensive comparison of multiple approaches to solve this issue:




Figure 1: Vocabulary Mismatch Problem: Traditional exact-match systems (TF-IDF) fail when query terms differ from document terms, even when they share the same meaning. Semantic models capture semantic equivalence through embeddings.

- **Baseline:** TF-IDF (Term Frequency-Inverse Document Frequency).
- **Static Embeddings:** Word2Vec trained on the corpus vs. Pre-trained BioWord2Vec.
- **Neural Models:** Deep Averaging Network (DAN) with triplet loss.
- **Transformer Models:** BioBERT (standard) and Sentence-BERT (optimized for semantic similarity).

Our contributions include: (1) a systematic evaluation framework using reusable functions for consistent model comparison, (2) detailed analysis of why certain approaches fail despite their theoretical advantages, and (3) insights into the trade-offs between model complexity and data availability in specialized domains.

The complete source code, notebooks, and experimental results are available at: https://github.com/Lounes-KEBDI/medical_information_retrieval.git

2 Related Work

2.1 TREC-COVID Challenge and Previous Approaches

The TREC-COVID challenge [8] was established in response to the urgent need for effective information retrieval systems during the COVID-19 pandemic. Conducted over five rounds from April to July 2020, this shared task attracted participation from 92 unique teams and generated 556 individual submissions. The challenge aimed to construct a pandemic search test collection and evaluate IR methods specifically for COVID-19 scientific literature. The dataset consists of 50 topics (queries) representing real-world information needs, with topics added progressively to target emerging research

questions as the pandemic evolved. This comprehensive evaluation framework has become a standard benchmark for biomedical information retrieval systems.

Previous work on TREC-COVID has explored various approaches. Top-performing systems in the challenge typically combined multiple techniques: lexical methods (BM25), neural reranking with BERT-based models, query expansion, and hybrid architectures. Many successful submissions leveraged domain-specific pre-trained models (BioBERT, PubMedBERT) with task-specific fine-tuning on retrieval pairs. However, most studies focused on achieving high performance through ensemble methods and complex pipelines, with limited analysis of why individual approaches succeed or fail. Few works have provided systematic comparisons across the full spectrum from simple baselines (TF-IDF) to state-of-the-art transformers, or investigated the trade-offs between model complexity and data requirements in this low-resource setting.

2.2 Evolution of Information Retrieval Methods

Information retrieval has evolved from lexical matching to semantic understanding. The Vector Space Model, exemplified by TF-IDF, remains a strong baseline but lacks semantic depth. It relies on exact word matching and term frequency statistics, making it vulnerable to vocabulary mismatch problems.

The introduction of Word Embeddings, such as Word2Vec [6], allowed for vector representations where similar words are close in space. These static embeddings capture semantic relationships but lack contextual understanding. In the medical domain, models pre-trained on biomedical corpora, such as BioBERT [5], have shown that domain adaptation is crucial for performance. However, BioBERT inherits the vocabulary of general BERT, causing medical terms to be fragmented into subword tokens. PubMedBERT [2] addresses this by training from scratch on biomedical corpora, constructing a domain-specific vocabulary that avoids term fragmentation and achieves superior performance on biomedical benchmarks.

Deep Averaging Networks (DAN) [3] combine word embeddings with neural architectures, averaging word vectors and passing them through hidden layers. While effective for classification tasks, their performance in retrieval depends heavily on training data availability.

More recently, Reimers et al. introduced Sentence-BERT [7], optimizing BERT architectures specifically for semantic similarity tasks using siamese networks. This approach addresses a critical limitation of standard BERT: the [CLS] token embedding is not optimized for cosine similarity in dense retrieval scenarios. Subsequent work has extended this paradigm to biomedical domains, with models like MedCPT [4] leveraging contrastive learning on large-scale search logs to align query and document embeddings. Hybrid approaches combining lexical (BM25) and semantic (dense) signals, such as SPLADE [1], have also shown promise by learning sparse lexical representations that maintain interpretability while benefiting from contextual understanding.

2.3 Our Contributions

This work addresses several gaps in previous TREC-COVID studies. First, we provide a **systematic evaluation framework** using

reusable, model-agnostic functions that ensure fair comparison across all approaches—from simple TF-IDF to complex transformers. This framework eliminates implementation biases and allows us to isolate the impact of model architecture choices.

Second, we conduct a **comprehensive comparative analysis** spanning the full spectrum of IR methods: lexical baselines (TF-IDF), static embeddings (Word2Vec variants), neural architectures (DAN), and transformer models (BioBERT, Sentence-BERT). Unlike previous work that focused on ensemble methods, we evaluate each approach independently to understand their individual strengths and limitations.

Third, we provide **detailed failure analysis** explaining why certain theoretically promising approaches (e.g., BioBERT, fine-tuned models) underperform despite their advantages. Our analysis reveals critical insights: (1) domain expertise alone is insufficient without task-specific optimization, (2) model complexity requires adequate training data, and (3) simple approaches can outperform complex architectures in low-resource scenarios.

Finally, we investigate the **trade-offs between model complexity and data availability** in specialized domains. Our experiments demonstrate that trainable models (DAN, fine-tuned Word2Vec) suffer from overfitting with limited queries (50), while pre-trained models achieve better generalization. These findings provide practical guidance for IR system design in resource-constrained biomedical settings.

3 Methodology

3.1 Dataset

We utilized the TREC-COVID dataset [9], accessed through the BEIR framework, which consists of:

- **Corpus:** 171,332 scientific articles from the CORD-19 collection, each containing a title and text (abstract or full content).
- **Queries:** 50 distinct topics in natural language, representing real-world information needs.
- **Qrels:** Expert relevance judgments used as Ground Truth, with relevance scores on a -1–2 scale (-1: Document not evaluated, 0: not relevant, 1: partially relevant, 2: highly relevant).

The dataset presents significant challenges: documents use technical medical terminology, while queries often employ lay language. This vocabulary mismatch is the primary obstacle our models must overcome.

3.2 Evaluation Framework

To ensure consistent and fair comparison across all models, we implemented a reusable evaluation framework in `functions.py`. This framework provides:

Figure 2 illustrates the general information retrieval pipeline used across all models.

To ensure consistent and fair comparison across all models, we implemented a reusable evaluation framework in `functions.py`. This framework provides:

1. Model-Agnostic Search Function: The `search()` function accepts any vectorizer/model and performs top-k retrieval using




Figure 2: Information Retrieval Pipeline: All models follow the same pipeline: (1) vectorize query and documents using model-specific methods, (2) compute cosine similarity, (3) rank documents, (4) retrieve top-K results, and (5) evaluate against expert judgments (Qrels).

cosine similarity. This abstraction allows us to evaluate TF-IDF, Word2Vec, and transformer models with identical search logic.

2. Standardized Evaluation: The `evaluate_model()` function computes NDCG@10 for each query and aggregates statistics (mean, min, max, standard deviation) across all 50 queries. This ensures that performance differences reflect model capabilities rather than implementation variations.

3. Ground Truth Comparison: The framework compares model predictions against expert judgments (Qrels) using scikit-learn’s `ndcg_score` implementation, ensuring metric consistency.

3.3 Models Implemented

3.3.1 TF-IDF (Baseline). Implementation: We used scikit-learn’s `TfidfVectorizer` with default parameters. Documents and queries are represented as sparse TF-IDF vectors, and retrieval is performed via cosine similarity.

Rationale: TF-IDF serves as our lower-bound reference. It captures term importance through frequency statistics but cannot handle semantic relationships or vocabulary mismatch.

3.3.2 Word2Vec Variants. 1. W2V-General (Trained from Scratch): We trained a Word2Vec model from scratch using the entire TREC-COVID corpus. **Training data:** All 171,332 documents from the corpus were used. For each document, we combined the title and text, tokenized using NLTK’s `word_tokenize`, removed stopwords, and filtered to keep only alphabetic tokens. This produced a training corpus of 171,332 tokenized documents with an average length of 94 words per document. **Technical implementation:** We used Gensim’s Word2Vec with the following parameters: 200-dimensional embeddings, window size 5 (context window of 5 words on each side), minimum word count 2 (words appearing less than twice are ignored), and trained for 5 epochs (5 complete passes over the corpus). The model learned embeddings for 105,771 unique words present in the TREC-COVID corpus. **Important note:** This model is trained unsupervised on the document corpus only; the 50 test queries are not used during training and are only used for evaluation.

2. BioWord2Vec (Pre-trained): We used a pre-trained model (BioWordVec_PubMed_MIMICIII_d200) trained on 1.5 million biomedical words from PubMed and MIMIC-III. This model provides domain-specific embeddings without requiring any training on our corpus. It was used directly for inference without modification. **3. BioWord2Vec Fine-tuned:** We fine-tuned the pre-trained BioWord2Vec model on the TREC-COVID corpus. **Training data:** All 171,332 documents from the TREC-COVID corpus were used, processed identically to the W2V-General model (title + text, tokenization, stopword removal). **Technical implementation:** We created a new Word2Vec model with the same architecture as the pre-trained model (200 dimensions, window size 5). We then built a vocabulary from the TREC-COVID documents, resulting in 81,225 unique words. For words present in both vocabularies (69,868 common words), we initialized their embeddings with the pre-trained weights, preserving the biomedical knowledge learned from PubMed/MIMIC-III. For words unique to TREC-COVID (11,357 new words), we initialized embeddings randomly and learned them during fine-tuning. The model was then trained for 5 epochs on the TREC-COVID documents, allowing it to adapt the embeddings to corpus-specific terminology while retaining general biomedical knowledge. **Important note:** As with W2V-General, only the document corpus is used for training; queries are not used during fine-tuning.

Document Representation: For all Word2Vec variants, documents are represented by averaging word embeddings. We combine title and text, tokenize using NLTK, remove stopwords, and compute the mean embedding vector. Queries are processed identically during inference (not during training).

3.3.3 Deep Averaging Network (DAN). Architecture: We implemented a bi-encoder architecture where queries and documents share the same encoder. The encoder consists of:

- Input layer: BioWord2Vec embeddings (200 dimensions)
- Hidden layer: Fully connected layer (300 dimensions) with ReLU activation

- Output layer: Document/query embedding (200 dimensions)

Figure 3 illustrates the DAN architecture and encoding process.




Figure 3: DAN Architecture: The encoder processes text through tokenization, BioWord2Vec embeddings, averaging, and a Multi-Layer Perceptron (MLP) to produce optimized 200-dimensional embeddings. The same encoder is used for both queries and documents (bi-encoder architecture).

Training: We used triplet margin loss with margin=0.5. For each query, we created triplets (query, positive document, negative document) based on Qrels relevance scores. Positive documents have relevance ≥ 1 , negative documents have relevance = 0. The model was trained for 5 epochs with Adam optimizer (learning rate = 0.001).

Challenges: With only 50 queries and limited positive examples per query, creating diverse triplets was challenging. We used random negative sampling, which may not provide optimal hard negatives.

3.3.4 Transformer Models. 1. BioBERT (Standard): We used the dmis-lab/biobert-base-cased-v1.1 model from Hugging Face. Documents and queries are tokenized with a maximum length of 512

tokens. We extract the [CLS] token embedding (768 dimensions) as the document/query representation. No fine-tuning was performed for the retrieval task.

2. Sentence-BERT (S-PubMedBert): We used pritamdeka/S-PubMedBert-MS-MARCO, a Siamese BERT network fine-tuned on MS-MARCO (general IR) and PubMed (biomedical domain). This model is specifically optimized for cosine similarity in semantic search. Documents and queries are encoded using `model.encode()`, which returns normalized embeddings (768 dimensions) ready for cosine similarity computation.

Key Difference: Sentence-BERT embeddings are optimized for semantic similarity through contrastive learning on pairs of similar/dissimilar texts. BioBERT’s [CLS] token is trained for classification tasks, not similarity search.

3.4 Evaluation Metric

All models were evaluated using **NDCG@10** (Normalized Discounted Cumulative Gain at 10), which prioritizes highly relevant documents at the top of the ranking list. NDCG accounts for both relevance and ranking position, making it ideal for IR evaluation. We report mean NDCG@10 across all 50 queries, along with standard deviation, minimum, and maximum scores.

4 Experiments and Results

4.1 Experimental Setup

All experiments were conducted on a standard CPU environment. For transformer models, we processed documents in batches to manage memory efficiently. Pre-computed document vectors were saved to disk (`models/`) to enable rapid query evaluation without recomputing embeddings.

Results for each model were saved in JSON format (`results/`) containing:

- Mean, min, max, and standard deviation of NDCG@10
- Individual NDCG@10 scores for each of the 50 queries
- Model-specific metadata (embedding dimensions, training epochs, etc.)

4.2 Performance Comparison

The comparative results of all implemented models are summarized in Table 1.

Table 1: Performance Comparison (Mean NDCG@10 across 50 queries)

Model	Mean	Std	Min	Max
TF-IDF (Baseline)	0.5868	0.3350	0.0000	1.0000
Word2Vec (Corpus)	0.6099	0.3301	0.0000	1.0000
BioWord2Vec (Pre-trained)	0.7366	0.3029	0.0000	1.0000
BioWord2Vec (Fine-tuned)	0.6259	0.3263	0.0000	1.0000
DAN	0.6630	0.2696	0.0000	1.0000
BioBERT (Standard)	0.2886	0.3454	0.0000	1.0000
Sentence-BERT	0.8171	0.2305	0.0000	0.9986

4.3 Detailed Analysis

1. TF-IDF Baseline: Achieved 0.5868 mean NDCG@10 with high variance (std=0.3350), indicating inconsistent performance across queries. Some queries achieved perfect scores (max=1.0) when exact term matching was sufficient, while others failed completely (min=0.0) due to vocabulary mismatch.

2. Word2Vec (Corpus-trained): Modest improvement (+3.9%) over TF-IDF. The model learned some semantic relationships within the TREC-COVID corpus but lacked general biomedical knowledge. Performance variance remained high (std=0.3301).

3. BioWord2Vec (Pre-trained): Significant improvement (+25.5% vs baseline, +20.8% vs corpus-trained Word2Vec). The pre-trained model's biomedical vocabulary (1.5M words from PubMed/MIMIC-III) effectively captured domain-specific semantics. Lower variance (std=0.3029) indicates more consistent performance.

4. BioWord2Vec (Fine-tuned): Surprisingly, fine-tuning degraded performance (-15.0% vs pre-trained version). Despite learning 11,357 new TREC-COVID-specific words, the model likely overfitted to the training corpus, losing the generalizability of pre-trained embeddings.

5. DAN: Achieved 0.6630 NDCG@10 (+13.0% vs baseline) but underperformed compared to simple BioWord2Vec averaging (-10.0%). The neural architecture's complexity required more training data than available (50 queries). Lower variance (std=0.2696) suggests the model learned some general patterns but lacked query-specific adaptation.

6. BioBERT (Standard): Catastrophic failure (0.2886 NDCG@10, -50.8% vs baseline). The [CLS] token embedding, designed for classification, is not optimized for cosine similarity in dense retrieval. Without task-specific fine-tuning, BioBERT embeddings are misaligned for semantic search.

7. Sentence-BERT: Best performance (0.8171 NDCG@10, +39.2% vs baseline, +10.9% vs BioWord2Vec). The model's optimization for semantic similarity, combined with biomedical domain pre-training, effectively bridges the vocabulary gap. Lowest variance (std=0.2305) indicates robust and consistent performance across diverse queries.

5 Discussion

5.1 Success of Semantic Search

The results confirm that semantic models generally outperform keyword matching. **Sentence-BERT** achieved the highest score (0.8171), demonstrating the power of contextual embeddings optimized for cosine similarity. It effectively bridges the vocabulary gap between lay queries and medical texts through:

- Domain-specific pre-training on PubMed
- Task-specific optimization for semantic similarity (MS-MARCO)
- Normalized embeddings designed for cosine similarity

5.2 The Importance of Pre-training

BioWord2Vec (0.7366) significantly outperformed the **Word2Vec** model trained only on the specific corpus (0.6099). This highlights that for specialized domains, transferring knowledge from massive

external datasets (PubMed, MIMIC-III) is more effective than training on a smaller, specific dataset. The pre-trained model's 1.5M word vocabulary captures biomedical relationships that would require orders of magnitude more data to learn from scratch.

5.3 Failures and Challenges

5.3.1 BioBERT Failure. The standard BioBERT model performed poorly (0.2886). This failure is attributed to several factors:

- **Task Mismatch:** The [CLS] token is trained for classification, not similarity search. Its embedding space is not optimized for cosine similarity between queries and documents.
- **No Fine-tuning:** Without task-specific fine-tuning on retrieval pairs, BioBERT embeddings lack the alignment necessary for effective dense retrieval.
- **Embedding Normalization:** Raw BERT embeddings are not normalized, making cosine similarity less meaningful compared to models explicitly trained for similarity.

This result highlights a critical lesson: *domain expertise alone is insufficient; task-specific optimization is essential.*

5.3.2 Overfitting in Trainable Models. Both the Fine-tuned BioWord2Vec (0.6259) and the DAN model (0.6630) showed lower performance than the pre-trained BioWord2Vec. With only 50 queries available for training/validation, these models likely suffered from overfitting, learning noise rather than generalizable patterns.

Fine-tuned BioWord2Vec: Despite learning 11,357 new TREC-COVID-specific words, the model's performance degraded. Possible explanations:

- Overfitting to corpus-specific patterns that don't generalize
- Loss of pre-trained knowledge through excessive fine-tuning (5 epochs)
- Insufficient regularization or early stopping

DAN Model: The neural architecture's capacity exceeded the available training data. With only 50 queries and limited positive examples per query, the model could not learn robust query-document relationships. The triplet loss formulation may have also suffered from poor negative sampling strategies.

6 Conclusion and Future Work

This study evaluated multiple IR strategies for the TREC-COVID challenge, ranging from classical lexical methods to state-of-the-art neural architectures. We demonstrated that **Sentence-BERT** is the most effective approach (NDCG@10: 0.8171), leveraging both domain-specific pre-training and an architecture optimized for semantic similarity.

6.1 Key Findings

- (1) **Domain-specific pre-training is crucial:** BioWord2Vec (+25.5%) and Sentence-BERT (+39.2%) significantly outperform generic models.
- (2) **Task-specific optimization matters:** Sentence-BERT's similarity-focused training outperforms BioBERT's classification-focused embeddings by 183%.

- (3) **Model complexity requires data:** Trainable models (DAN, fine-tuned W2V) underperform simpler approaches when training data is limited (50 queries).
- (4) **Simple can be effective:** BioWord2Vec averaging (0.7366) outperforms complex neural architectures (DAN: 0.6630) in low-data scenarios.

6.2 Future Work

A major limitation observed was the tendency of trainable models (DAN, Fine-tuned W2V) to overfit due to the small number of queries (50). Future research should focus on:

1. Data Augmentation: Techniques such as using Large Language Models (LLMs) to synthetically generate variations of queries and pseudo-relevance judgments could expand the training set. This would allow for robust training of deeper architectures without the risk of overfitting observed in this study.

2. Hard Negative Mining: For triplet-based training (DAN), implementing sophisticated negative sampling strategies could improve learning efficiency. Hard negatives (documents that are similar but not relevant) provide more informative training signals than random negatives.

3. Cross-Domain Transfer: Exploring transfer learning from other biomedical IR datasets (e.g., TREC Genomics, BioASQ) could provide additional training signals while maintaining domain relevance.

4. Hybrid Approaches: Combining lexical (TF-IDF) and semantic (Sentence-BERT) signals through learned fusion mechanisms could capture both exact matches and semantic relationships.

5. Query Expansion: Pre-processing queries with medical term expansion (e.g., mapping “loss of smell” to “anosmia”) could improve performance for models that struggle with vocabulary mismatch.

6. Fine-tuning Strategies: Investigating optimal fine-tuning protocols for BioBERT on retrieval tasks, including contrastive learning objectives and appropriate learning rates, could bridge the performance gap with Sentence-BERT.

6.3 Broader Implications

This study highlights the importance of matching model architecture and training objectives to the target task. While BioBERT excels at biomedical text understanding, its standard configuration fails for retrieval without task-specific adaptation. This underscores the need for specialized models (like Sentence-BERT) or careful fine-tuning when applying general-purpose architectures to specific IR scenarios.

References

- [1] Thibault Formal, Benjamin Piwowarski, and Stéphane Clinchant. 2021. SPLADE: Sparse Lexical and Expansion Model for First Stage Ranking. *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval* (2021), 2288–2292.
- [2] Yu Gu, Robert Tinn, Hao Cheng, Michael Lucas, Naoto Usuyama, Xiaodong Liu, Tristan Naumann, Jianfeng Gao, and Hoifung Poon. 2020. Domain-Specific Language Model Pretraining for Biomedical Natural Language Processing. *arXiv preprint arXiv:2007.15779* (2020).
- [3] Mohit Iyyer, Varun Manjunatha, Jordan Boyd-Graber, and Hal Daumé III. 2015. Deep unordered composition rivals syntactic methods for text classification. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics*. 1681–1691.
- [4] Qiao Jin, Bhuwan Dhingra, William W Cohen, and Xinghua Lu. 2023. MedCPT: Contrastive Pre-trained Transformers with Large-scale PubMed Search Logs for Zero-shot Biomedical Information Retrieval. *Bioinformatics* 39, 11 (2023), btad652.
- [5] Jinhyuk Lee, Wonjin Yoon, Sungdong Kim, Donghyeon Kim, Sunkyu Kim, Chan Ho So, and Jaewoo Kang. 2020. BioBERT: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics* 36, 4 (2020), 1234–1240.
- [6] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013. Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781* (2013).
- [7] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*. 3982–3992.
- [8] Kirk Roberts, Tasmeer Alam, Steven Bedrick, Dina Demner-Fushman, Kyle Lo, Ian Soboroff, Ellen Voorhees, Lucy Lu Wang, and William R Hersh. 2021. Searching for Scientific Evidence in a Pandemic: An Overview of TREC-COVID. *arXiv preprint arXiv:2104.09632* (2021).
- [9] Nandan Thakur, Nils Reimers, Andreas Rücklé, Abhishek Srivastava, and Iryna Gurevych. 2021. BEIR: A Heterogeneous Benchmark for Zero-shot Evaluation of Information Retrieval Models. *arXiv preprint arXiv:2104.08663* (2021).
- [10] Ellen Voorhees, Tasmeer Alam, Steven Bedrick, Dina Demner-Fushman, William R Hersh, Jimmy Lin, Kyle Lo, Kirk Roberts, Ian Soboroff, Lucy Lu Wang, et al. 2021. TREC-COVID: constructing a pandemic information retrieval test collection. *ACM SIGIR Forum* 54, 1 (2021), 1–12.